

RUL Prediction as Classification

RUL Prediction as Classification

RUL-based maintenance can also be tackled **using a classifier**

- We build a classifier to determine whether a failure will occur in θ steps
- We stop as soon as the classifier outputs (say) a 0, i.e.

$$f_{\theta}(x, \lambda) = 0$$

- f is the classifier, with parameter vector λ
- θ is the horizon for detecting a failure

In a sense, we are trying to learn **directly a maintenance policy**

- The policy is the form "stop θ units before a failure"
- The classifier tries to learn it

Classifier Architecture

We can therefore immediately define our classifier architecture:

```
In [2]: def build_classifier(hidden):  
    input_shape = (len(dt_in), )  
    model_in = keras.Input(shape=input_shape, dtype='float32')  
    x = model_in  
    for h in hidden:  
        x = layers.Dense(h, activation='relu')(x)  
    model_out = layers.Dense(1, activation='sigmoid')(x)  
    model = keras.Model(model_in, model_out)  
    return model
```

- Like in the regression case, we use a Multilayer Perceptron
- The only difference is the use of a sigmoid activation in the output layer
- For `hidden = []` we get Logistic Regression

Training

Before training, we need to define the classes

In turn, this requires to define the detection horizon θ :

```
In [3]: class_thr = 15
        tr_lbl = (tr['rul'] >= class_thr)
        ts_lbl = (ts['rul'] >= class_thr)
```

- The class is "1" if a failure is more than θ steps away
- The class is "0" otherwise

Classification problems **tend to be easier** than regression problems

- On the other hand, learning the whole policy
- ...May be trickier than just estimating the RUL

Training

Let's start by training the simplest possible model

```
In [4]: nn1 = build_classifier(hidden=[])
nn1.compile(optimizer='Adam', loss='binary_crossentropy')
cb = [callbacks.EarlyStopping(patience=10, restore_best_weights=True)]
history1 = nn1.fit(tr_s[dt_in], tr_lbl, validation_split=0.2,
                  callbacks=cb,
                  batch_size=32, epochs=20, verbose=1)
```

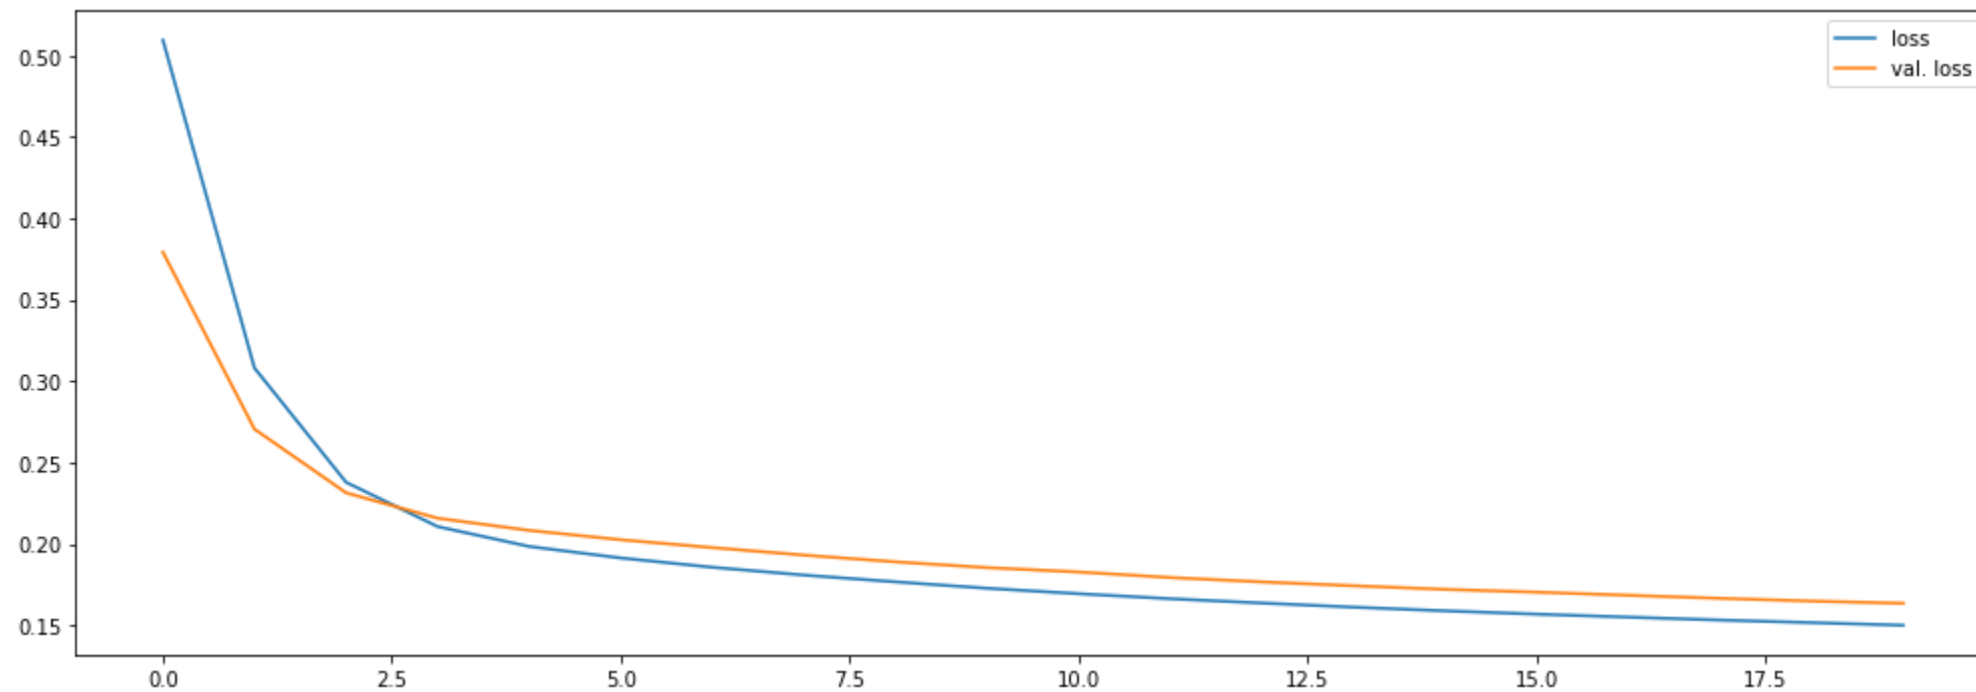
```
Epoch 1/20
1135/1135 [=====] - 1s 636us/step - loss: 0.5097 - val_loss: 0.3793
Epoch 2/20
1135/1135 [=====] - 1s 552us/step - loss: 0.3081 - val_loss: 0.2705
Epoch 3/20
1135/1135 [=====] - 1s 549us/step - loss: 0.2378 - val_loss: 0.2313
Epoch 4/20
1135/1135 [=====] - 1s 556us/step - loss: 0.2106 - val_loss: 0.2157
Epoch 5/20
1135/1135 [=====] - 1s 551us/step - loss: 0.1983 - val_loss: 0.2082
Epoch 6/20
1135/1135 [=====] - 1s 554us/step - loss: 0.1913 - val_loss: 0.2025
Epoch 7/20
1135/1135 [=====] - 1s 556us/step - loss: 0.1856 - val_loss: 0.1976
Epoch 8/20
1135/1135 [=====] - 1s 552us/step - loss: 0.1808 - val_loss: 0.1930
Epoch 9/20
1135/1135 [=====] - 1s 550us/step - loss: 0.1766 - val_loss: 0.1889
Epoch 10/20
1135/1135 [=====] - 1s 548us/step - loss: 0.1726 - val_loss: 0.1852
```

Training

Here's the loss evolution over time and its final value

```
In [7]: cmapss.plot_training_history(history1, figsize=figsize)
tr1, vl1 = history1.history["loss"][-1], np.min(history1.history["val_loss"])
print(f'Final loss: {tr1:.4f} (training), {vl1:.4f} (validation)')
```

Final loss: 0.1499 (training), 0.1634 (validation)



Training

Let's try with a deeper model

```
In [8]: nn2 = build_classifier(hidden=[32, 32])
nn2.compile(optimizer='Adam', loss='binary_crossentropy')
history2 = nn2.fit(tr_s[dt_in], tr_lbl, validation_split=0.2,
                  callbacks=cb,
                  batch_size=32, epochs=20, verbose=1)
```

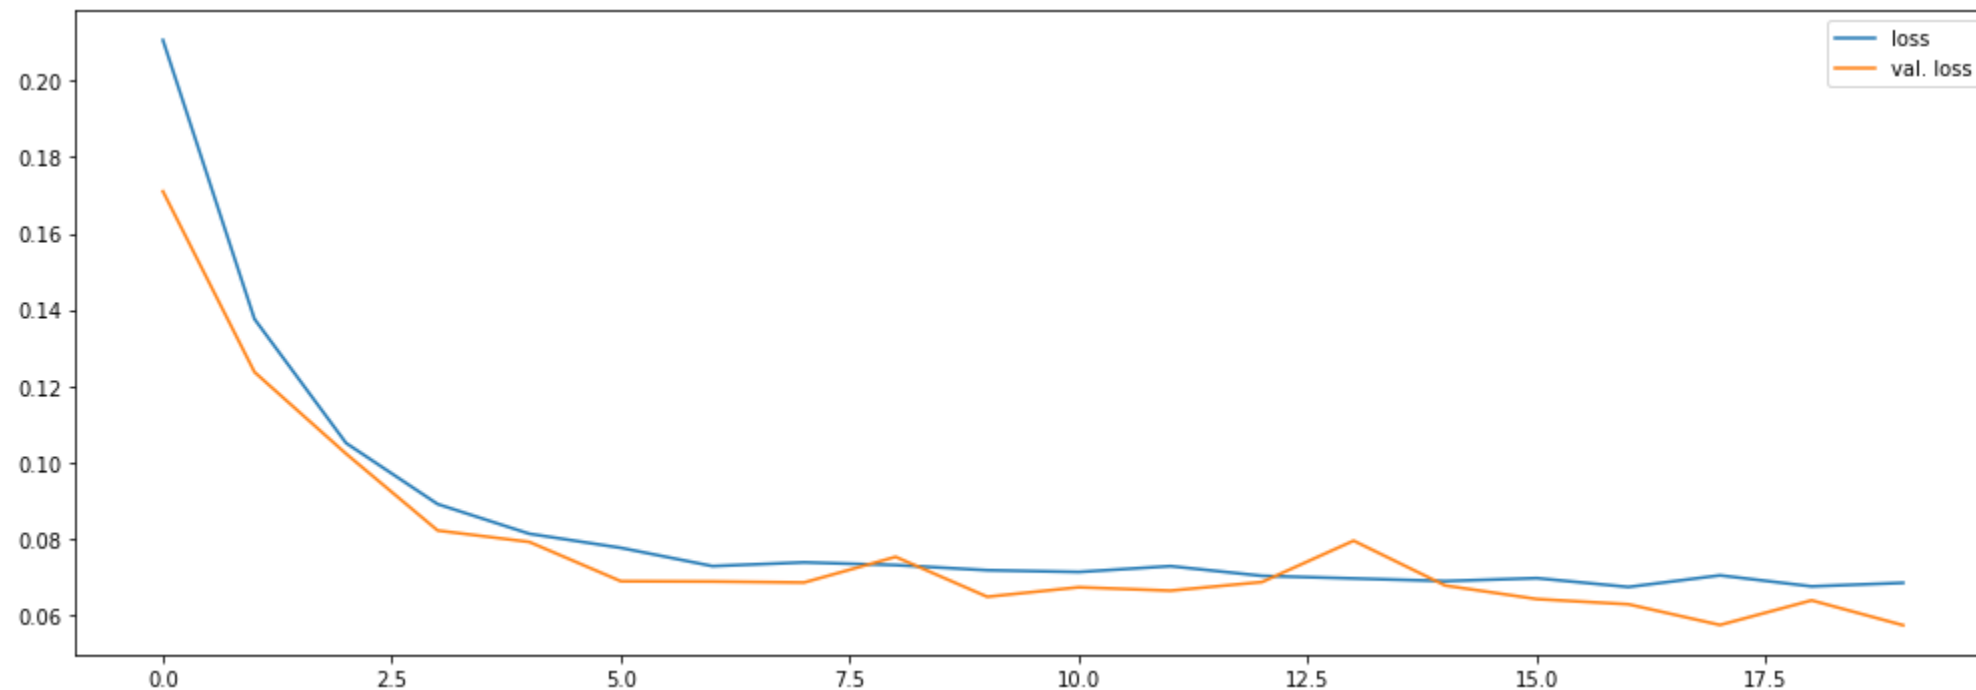
```
Epoch 1/20
1135/1135 [=====] - 1s 733us/step - loss: 0.2107 - val_loss: 0.1711
Epoch 2/20
1135/1135 [=====] - 1s 683us/step - loss: 0.1376 - val_loss: 0.1237
Epoch 3/20
1135/1135 [=====] - 1s 691us/step - loss: 0.1050 - val_loss: 0.1022
Epoch 4/20
1135/1135 [=====] - 1s 658us/step - loss: 0.0890 - val_loss: 0.0821
Epoch 5/20
1135/1135 [=====] - 1s 654us/step - loss: 0.0813 - val_loss: 0.0791
Epoch 6/20
1135/1135 [=====] - 1s 658us/step - loss: 0.0776 - val_loss: 0.0688
Epoch 7/20
1135/1135 [=====] - 1s 656us/step - loss: 0.0728 - val_loss: 0.0687
Epoch 8/20
1135/1135 [=====] - 1s 660us/step - loss: 0.0738 - val_loss: 0.0685
Epoch 9/20
1135/1135 [=====] - 1s 739us/step - loss: 0.0731 - val_loss: 0.0752
Epoch 10/20
1135/1135 [=====] - 1s 683us/step - loss: 0.0717 - val_loss: 0.0647
Epoch 11/20
```

Training

Here's the loss for the deeper model

```
In [9]: cmapss.plot_training_history(history2, figsize=figsize)
tr2, vl2 = history2.history["loss"][-1], np.min(history2.history["val_loss"])
print(f'Final loss: {tr2:.4f} (training), {vl2:.4f} (validation)')
```

Final loss: 0.0684 (training), 0.0573 (validation)

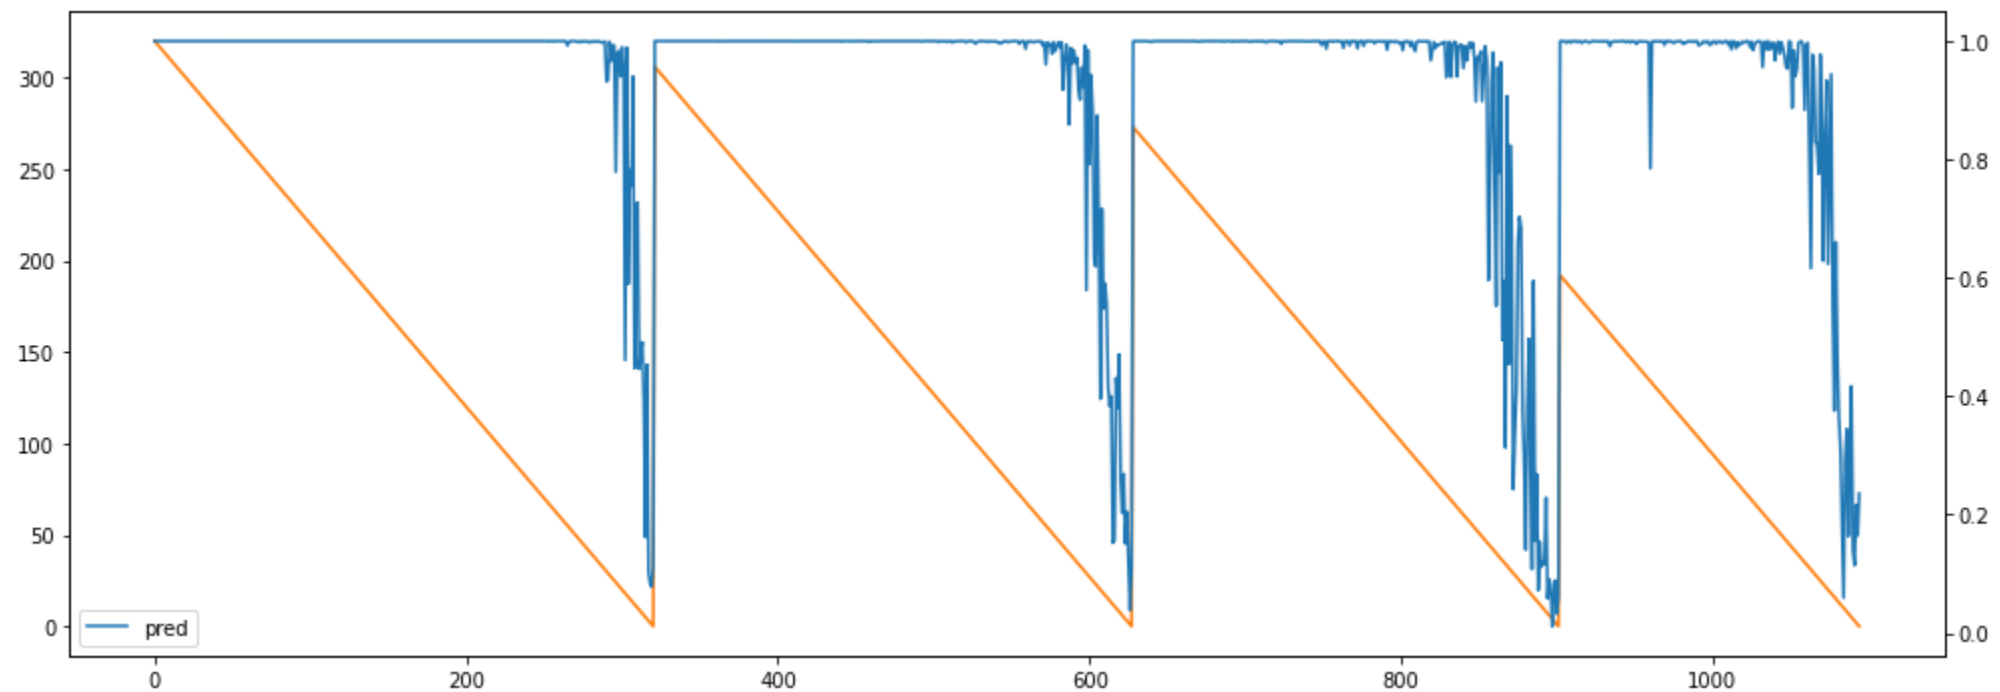


- Depth pays off big in this case

Predictions

Our predictions can be interpreted as probabilities (of advancing)

```
In [10]: tr_pred2_prob = nn2.predict(tr_s[dt_in]).ravel()  
stop = 1095  
cmapss.plot_rul(tr_pred2_prob[:stop], tr['rul'][:stop], same_scale=False, figsize=figsize)
```

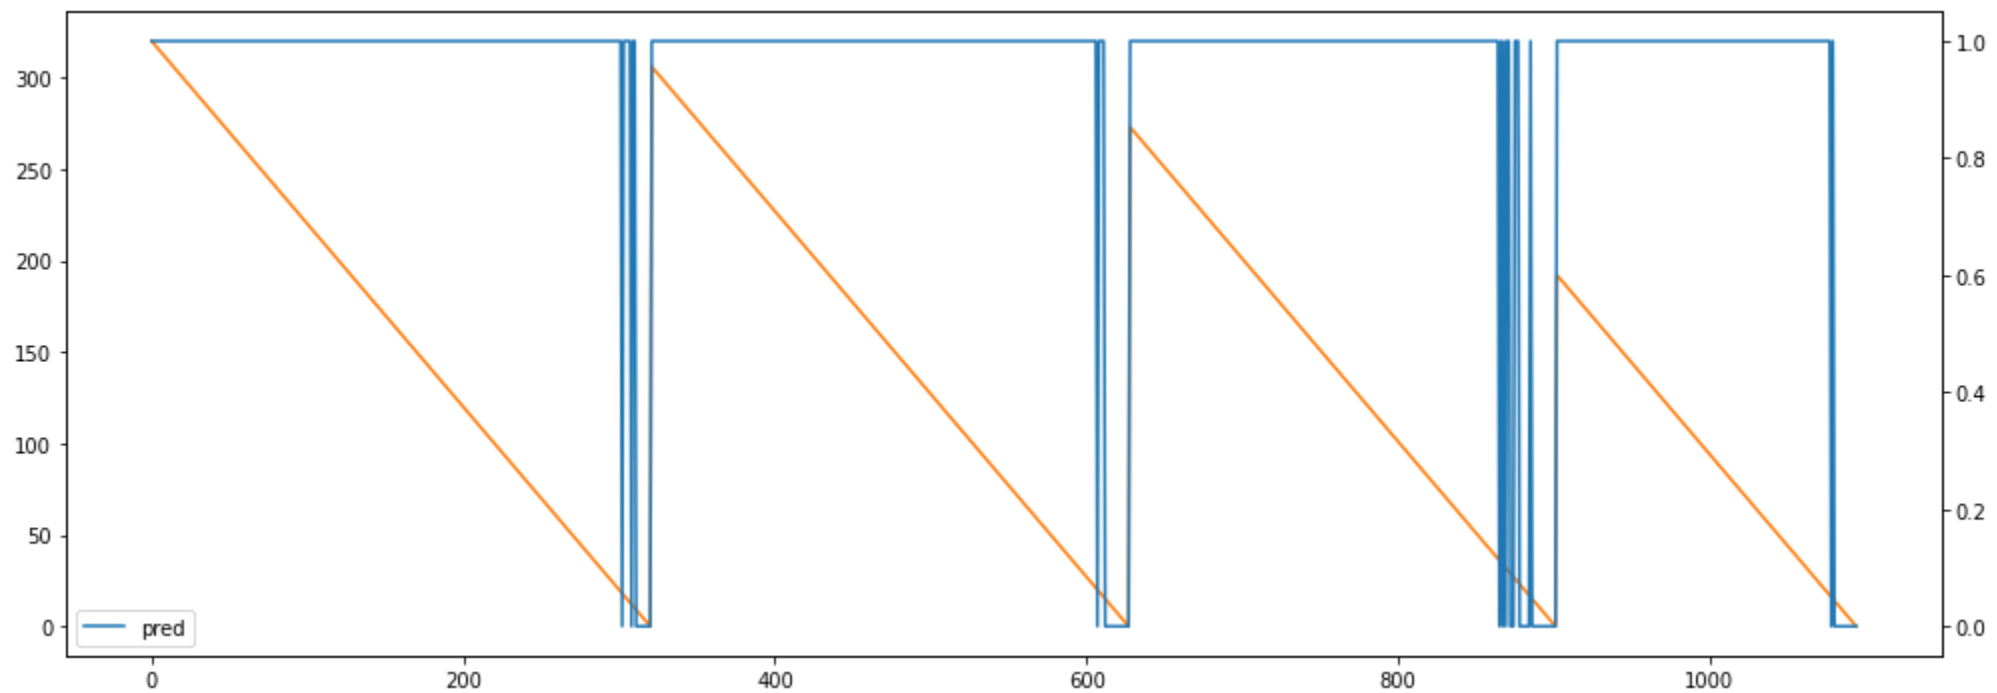


- The probability falls when closer to failures

Predictions

In practice, we need to convert the predictions to integers via rounding
...Unless we want to deal with one more threshold (in addition to θ)

```
In [11]: tr_pred2 = np.round(nn2.predict(tr_s[dt_in]).ravel())  
cmapss.plot_rul(tr_pred2[:stop], tr['rul'][:stop], same_scale=False, figsize=figsize)
```

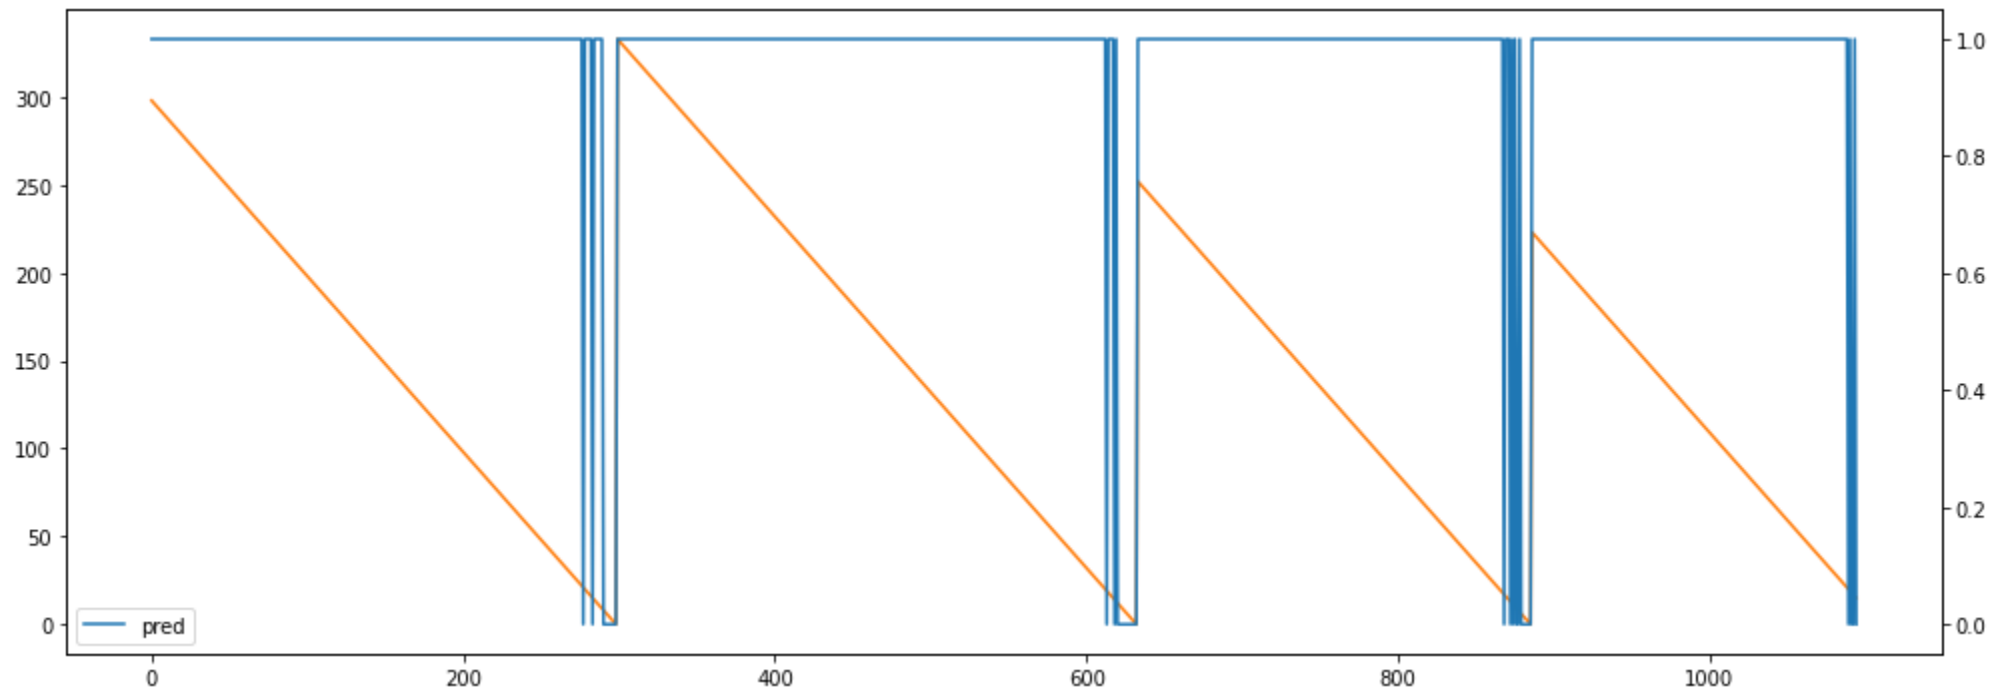


- Still, the behavior seems to be reasonable

Predictions

Let's see the behavior on the test set

```
In [12]: ts_pred2 = np.round(nn2.predict(ts_s[dt_in]).ravel())  
cmapss.plot_rul(ts_pred2[:stop], ts['rul'][:stop], same_scale=False, figsize=figsize)
```



- Apparently a decent degree of generalization

Evaluation

We can evaluate the classifier **directly**

...Because it defines the whole policy, with no need for additional calibration!

- On one hand this makes this stage of the process **simpler**
- ...On the other, this is (apparently) a missed opportunity

```
In [13]: tr_c2, tr_f2, tr_s2 = cmodel.cost(tr['machine'].values, tr_pred2, 0.5, return_margin=True)
         ts_c2, ts_f2, ts_s2 = cmodel.cost(ts['machine'].values, ts_pred2, 0.5, return_margin=True)
         print(f'Cost: {tr_c2} (training), {ts_c2} (test)')
         print(f'Avg. fails: {tr_f2/len(tr_mcn)} (training), {ts_f2/len(ts_mcn)} (test)')
         print(f'Avg. slack: {tr_s2/len(tr_mcn):.2f} (training), {ts_s2/len(ts_mcn):.2f} (test)')
```

```
Cost: -17221 (training), -6501 (test)
Avg. fails: 0.0 (training), 0.0 (test)
Avg. slack: 23.73 (training), 20.84 (test)
```

- Still pretty good results, but worse than the best regression approach
- We will investigate **two potential reasons**

Balancing a Dataset

Balancing a Dataset

By construction, our dataset is likely to be strongly unbalanced

```
In [14]: counts = tr_lbl.value_counts(normalize=True)
counts
```

```
Out[14]: True      0.938526
         False    0.061474
         Name: rul, dtype: float64
```

- In these kind of situation, SGD optimization may have convergence issues
- The gradient will push strongly in the direction of the overrepresented class

A common practice to address this issue is using class weights

- Typically, we use weights inversely proportional to the counts
- ...So as to counter-balance the effect

```
In [15]: class_weight = {0: 1/counts[0], 1: 1/counts[1]}
```

Class-Weights, Cross-Entropy, and Likelihood

When are class frequencies typically used as weights?

Let's say our classifier is trained via **weighted cross-entropy**:

$$\min_{\mathbf{y}} \left\{ -w_1 \sum_{\hat{y}_i=1} \log y_i - w_0 \sum_{\hat{y}_i=0} \log(1 - y_i) \right\}$$

- Where $\hat{\mathbf{y}}$ is the label vector and \mathbf{y} the vector of classifier outputs
- w_0 and w_1 are the weights for class 0 and 1

If we apply an exponential and we switch optimization direction we get:

$$\max_{\mathbf{y}} \left\{ e^{w_1 \sum_{\hat{y}_i=1} \log y_i} e^{w_0 \sum_{\hat{y}_i=0} \log(1-y_i)} \right\}$$

- The **optimal points** are the same as before

Class-Weights, Cross-Entropy, and Likelihood

We with algebraic manipulations we finally obtain:

$$\max_y \left\{ \prod_{\hat{y}_i=1} y_i^{w_1} \prod_{\hat{y}_i=0} (1 - y_i)^{w_0} \right\}$$

When training, we are **maximizing a likelihood expression!**

- To see that, just interpret y_i as the (estimated) probability of the class being 1

The weights act by multiplying the class probabilities

...Which is equivalent to adjusting their frequencies!

- E.g. from a mathematical point of view, setting w_0 to 2...
- ...Is equivalent to making a copy of each example having label 0

Hence, inverse class frequencies are often used as weights for dataset rebalancing

Training with Non-uniform Class Weights

Let's repeat the training process

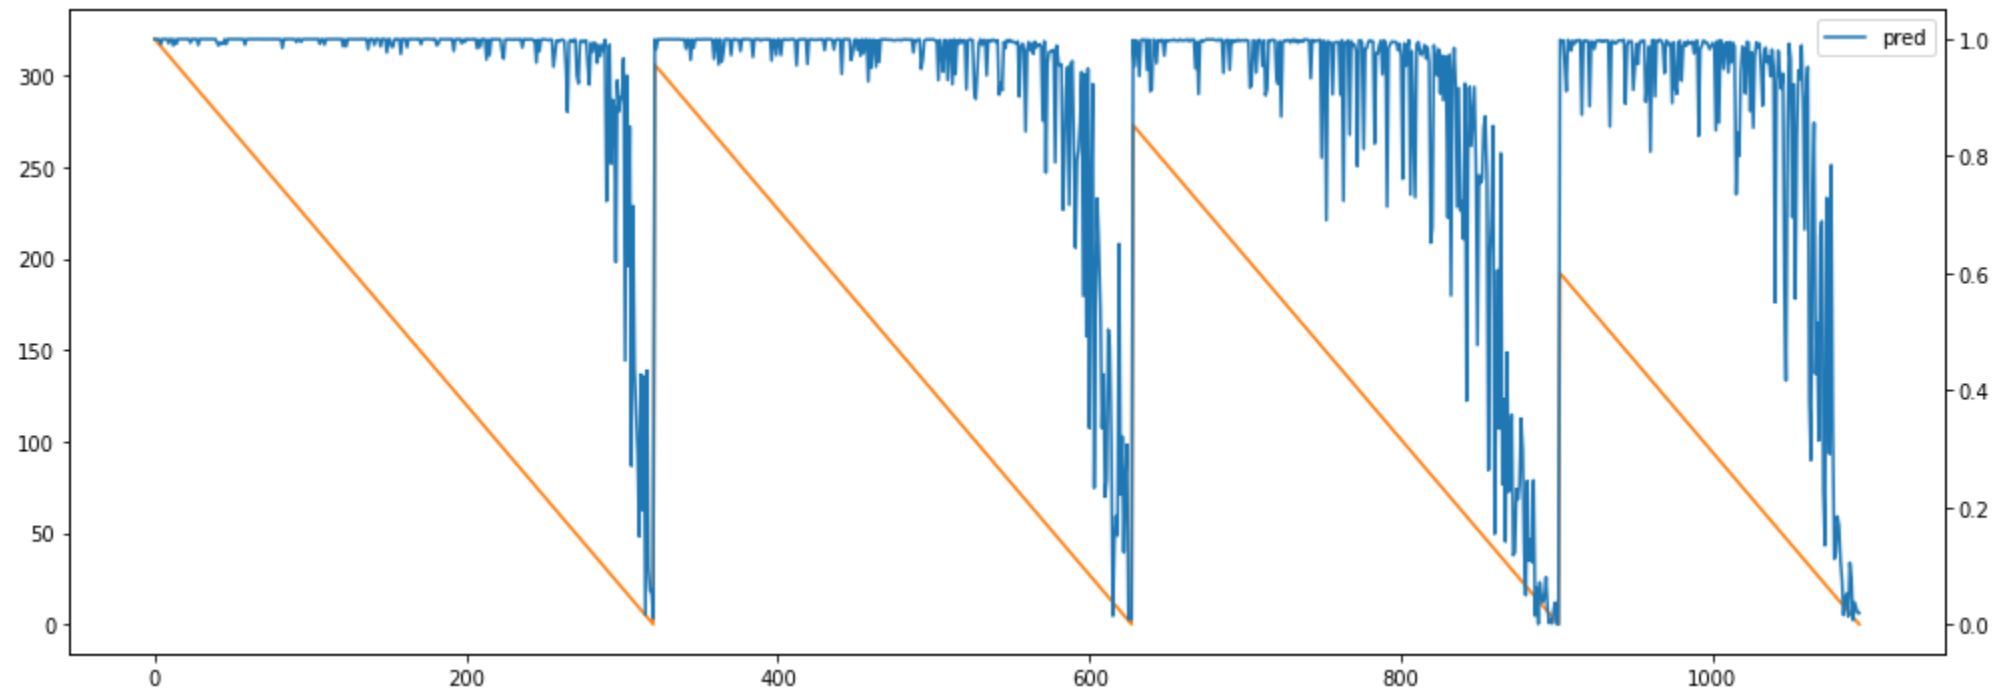
```
In [16]: nn3 = build_classifier(hidden=[32, 32])
nn3.compile(optimizer='Adam', loss='binary_crossentropy')
history3 = nn3.fit(tr_s[dt_in], tr_lbl, validation_split=0.2,
                  callbacks=cb, class_weight=class_weight,
                  batch_size=32, epochs=20, verbose=1)
```

```
Epoch 1/20
1135/1135 [=====] - 1s 773us/step - loss: 0.9785 - val_loss: 0.4846
Epoch 2/20
1135/1135 [=====] - 1s 699us/step - loss: 0.5663 - val_loss: 0.2182
Epoch 3/20
1135/1135 [=====] - 1s 698us/step - loss: 0.4288 - val_loss: 0.2280
Epoch 4/20
1135/1135 [=====] - 1s 698us/step - loss: 0.3668 - val_loss: 0.1700
Epoch 5/20
1135/1135 [=====] - 1s 695us/step - loss: 0.3452 - val_loss: 0.0957
Epoch 6/20
1135/1135 [=====] - 1s 704us/step - loss: 0.3176 - val_loss: 0.1723
Epoch 7/20
1135/1135 [=====] - 1s 704us/step - loss: 0.3070 - val_loss: 0.1606
Epoch 8/20
1135/1135 [=====] - 1s 709us/step - loss: 0.2940 - val_loss: 0.3057
Epoch 9/20
1135/1135 [=====] - 1s 718us/step - loss: 0.2820 - val_loss: 0.1319
Epoch 10/20
1135/1135 [=====] - 1s 703us/step - loss: 0.2792 - val_loss: 0.1631
Epoch 11/20
```

Predictions

Let's check how the the raw predictions (probabilities) have changed

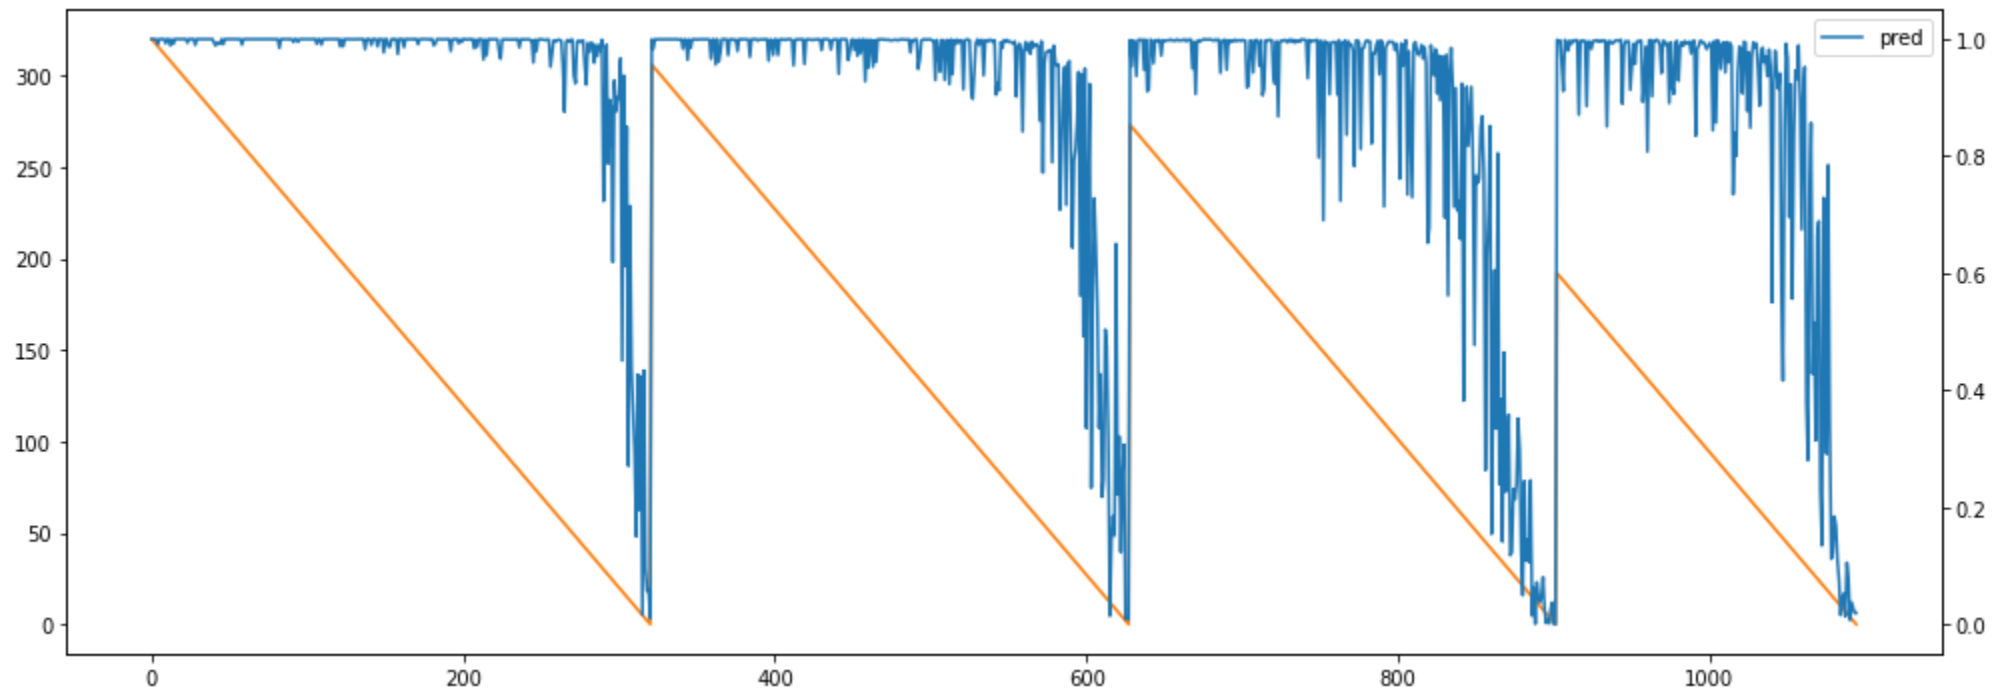
```
In [17]: tr_pred3_prob = nn3.predict(tr_s[dt_in]).ravel()  
cmapss.plot_rul(tr_pred3_prob[:stop], tr['rul'][:stop], same_scale=False, figsize=figsize)
```



Predictions

Let's check how the the raw predictions (probabilities) have changed

```
In [17]: tr_pred3_prob = nn3.predict(tr_s[dt_in]).ravel()  
cmapss.plot_rul(tr_pred3_prob[:stop], tr['rul'][:stop], same_scale=False, figsize=figsize)
```



- They have actually become worse: there are mistakes in the initial section

Evaluation

Let's check the model performance in terms of cost

```
In [19]: tr_pred3 = np.round(nn3.predict(tr_s[dt_in]).ravel())
ts_pred3 = np.round(nn3.predict(ts_s[dt_in]).ravel())
tr_c3, tr_f3, tr_s3 = cmodel.cost(tr['machine'].values, tr_pred3, 0.5, return_margin=True)
ts_c3, ts_f3, ts_s3 = cmodel.cost(ts['machine'].values, ts_pred3, 0.5, return_margin=True)
print(f'Cost: {tr_c3} (training), {ts_c3} (test)')
print(f'Avg. fails: {tr_f3/len(tr_mcn)} (training), {ts_f3/len(ts_mcn)} (test)')
print(f'Avg. slack: {tr_s3/len(tr_mcn):.2f} (training), {ts_s3/len(ts_mcn):.2f} (test)')
```

```
Cost: -15327 (training), -5871 (test)
Avg. fails: 0.0 (training), 0.0 (test)
Avg. slack: 34.12 (training), 31.44 (test)
```

As expected, the results are worse

- The model now treats the two classes more fairly
- ...But this increases the chance of an (undesirable) early stop

So, how shall we choose the balance?

Bayesian (Surrogate-Based) Optimization

Taking a Step Back

In the **regression** case, we are formally solving:

$$\min_{\theta} \sum_{k \in K} \text{cost}(f(x_k, \lambda^*), \theta)$$

$$\text{with: } \lambda^* = \operatorname{argmin}_{\lambda} L(f(x_k, \lambda), y_k)$$

- Where λ^* is the optimal parameter vector (i.e. the network weights)
- ...And L is the loss function (i.e. the MSE), and **cost** is our cost model
- θ is chosen so as to minimize the cost

This is a **bilevel optimization problem**

- However, since θ does not appear in L nor in f
- ...It can be **decomposed into two sequential subproblems**

Taking a Step Back

In the **classification** case, we are formally solving:

$$\min_{\theta} \sum_{k \in K} \text{cost}(f(x_k, \lambda^*), 0.5)$$

$$\text{with: } \lambda^* = \operatorname{argmin}_{\lambda} L(f(x_k, \lambda), \mathbf{1}_{y_k \geq \theta})$$

- In this case, we use a dummy threshold in the cost model (i.e. 0.5)
- ...And $\mathbf{1}_{y_k \geq \theta}$ is the indicator function of $y_k \geq \theta$ (i.e. our class labels)

Unlike the previous one, this problem cannot be decomposed

...Because θ appears in the loss function!

- This means we need to **optimize θ and λ at the same time**
- This is obviously complicated, but it's also an **opportunity**
- ...Since we can **adapt the training problem** based on the threshold

Taking a Step Back

We can try to tackle the optimization process in a hierarchical fashion

1. We search over the possible values of θ
2. For the given θ value, we compute $\mathbf{1}_{y_k \geq \theta}$ (i.e. the class labels)
3. For the given θ and $\mathbf{1}_{y_k \geq \theta}$, we compute λ^* (i.e. we train the model)
4. For the given θ , $\mathbf{1}_{y_k \geq \theta}$, and λ^* , we compute the cost

At the end of the process, we choose the configuration with the best cost

The method is reasonable, but slow

- We could reduce the runtime by warm-starting each training attempt
- ...And by searching over θ in a smart way

Bayesian Optimization

We will use an approach known as **Bayesian optimization**

- It's actually a family of **surrogate-based optimization** methods
- They are all designed to optimize **blackbox functions**
- I.e. functions with an unknown structure, that can only be evaluated

Formally, they address problems in the form:

$$\min_{x \in B} f(x)$$

- Where **B** is a box, i.e. a specification of bounds for each component of **x**

In our case:

- The decision variable **x** would be **θ**
- The function to be optimized would be the cost

Bayesian (or Surrogate-Based) Optimization

Internally, methods in this family store:

- A collection $\hat{\mathbf{x}}$ of evaluated points
- A surrogate-model \tilde{f} for f

The basic loop of the algorithm is then as follow:

Let \mathbf{x}^* be the current candidate optimum

- Based on the current \tilde{f} :
 - Determine whether there exists $\mathbf{x}' \in \mathcal{B}$ that may improve over \mathbf{x}^*
 - If there is no such point, return \mathbf{x}^*
 - Otherwise
 - Evaluate $f(\mathbf{x}')$ and possibly replace \mathbf{x}^*
 - Add \mathbf{x}' to $\hat{\mathbf{x}}$, adjust \tilde{f} , and repeat

Bayesian (or Surrogate-Based) Optimization

A few considerations:

In practice, the algorithm needs to balance **exploration** and **exploitation**

- We need to explore regions where we cannot make confident predictions
- ...But we need also to focus on regions with promising cost values

Different Bayesian optimization algorithms:

- Make use of different surrogate models
- Rely on different criteria for choosing \mathbf{x}'
- Strike different trade-offs in terms of number of (expensive) evaluations of f
- ...And the quality of the obtained solutions

For more information, see (e.g.) [this tutorial](#)

Our Chosen Optimizer

We will use the `bayesian-optimization` python module

- The solver is designed for maximization
- Decent documentation can be found on [its github page](#)
- It's not as fast as (e.g.) [RBFOpt](#), but easier to install and configure

The solver relies on **Gaussian Processes as a surrogate model:**

GPs are guaranteed to (approximately) touch every point in the training set

- ...And behave reasonably well even with a simple RBF kernel
- ...Which is in fact the one used in this case

GPs provide confidence intervals

...Which the algorithm uses to balance exploration and exploitation

- Regions with wide C.I. may be worth exploring
- ...And regions with high (upper) C.I. may contain high-quality solutions

Method Setup

First, we need to define the function to be maximized

This will be the (negative) cost, after retraining the classifier:

```
In [20]: nn4 = build_classifier(hidden=[32, 32])
nn4.compile(optimizer='Adam', loss='binary_crossentropy')
nn4.set_weights(nn2.get_weights())

bo_weights = {}
def classification_obj(class_thr, class0_weight):
    # Define new labels
    tr_lbl = (tr['rul'] >= class_thr)
    # Fit
    nn4.fit(tr_s[dt_in], tr_lbl, validation_split=0.2,
           callbacks=cb, class_weight={1:1, 0:class0_weight},
           batch_size=32, epochs=5, verbose=0)
    # Cost computation
    tr_pred = np.round(nn4.predict(tr_s[dt_in]).ravel())
    tr_cost = cmodel.cost(tr['machine'].values, tr_pred, 0.5)
    # Store the model weights
    bo_weights[class_thr, class0_weight] = nn4.get_weights()
    return -tr_cost
```

Method Setup

We rely on warm-starting to reduce the training time

We start from (the weights of) our deep MLP

```
nn4 = build_classifier(hidden=[32, 32])
nn4.compile(optimizer='Adam', loss='binary_crossentropy')
nn4.set_weights(nn2.get_weights())
```

...So that we can train for just a few epochs

```
nn4.fit(tr_s[dt_in], tr_lbl, validation_split=0.2,
        callbacks=cb, class_weight={1:1, 0:class0_weight},
        batch_size=32, epochs=5, verbose=0)
```

- Since the training process is stochastic...
- The optimization result will be stochastic as well

Method Setup

We will optimize **both θ and the class weights**

```
def classification_obj(class_thr, class0_weight):  
    ...  
    nn4.fit(tr_s[dt_in], tr_lbl, validation_split=0.2,  
           callbacks=cb, class_weight={1:1, 0:class0_weight},  
           batch_size=32, epochs=5, verbose=0)  
    ...
```

- We can afford it, since we have a much better optimizer than grid search

We need to store the weights of each model

...So that we can recover the optimal configuration:

```
bo_weights[class_thr, class0_weight] = nn4.get_weights()
```

Method Setup

We now define the bounding box

```
In [21]: box = {'class_thr': (1, 15), 'class0_weight': (1, 10)}
```

- We will consider stopping from 1 to 15 steps before a failure
- We will consider weights for class 0 from 1 to 10

Then we can configure the optimizer:

```
In [22]: from bayes_opt import BayesianOptimization
optimizer = BayesianOptimization(f=classification_obj,
                                pbounds=box, random_state=42, verbose=2)
```

- We specify the function to be maximized, the box
- ...But also a random seed and the logging level

Optimization

Finally, we run the optimization process

```
In [23]: optimizer.maximize(n_iter=10, init_points=3)
```

iter	target	class0...	class_thr
1	1.59e+04	4.371	14.31
2	1.569e+0	7.588	9.381
3	1.879e+0	2.404	3.184
4	1.742e+0	4.452	1.02
5	1.899e+0	1.0	5.556
6	-2.481e+0	1.0	1.0
7	1.881e+0	1.051	7.493
8	1.871e+0	1.678	6.5
9	1.779e+0	4.317	4.282
10	1.82e+04	7.043	2.12
11	1.717e+0	3.525	9.549
12	1.71e+04	1.0	11.22
13	1.654e+0	7.927	5.157

Retrieving the Results

We can access the best configuration via

```
In [24]: optimizer.max
```

```
Out[24]: {'target': 18994.0,  
          'params': {'class0_weight': 1.0000498300816418,  
                    'class_thr': 5.5561938101894865}}
```

Once we know the configuration, we can reconstruct the corresponding model:

```
In [25]: class_thr4 = optimizer.max['params']['class_thr']  
class0_weight4 = optimizer.max['params']['class0_weight']  
nn4.set_weights(bo_weights[class_thr4, class0_weight4])
```

- We do this by accessing the weights in the `bo_weights` dictionary

Evaluation

Finally, we can obtain the predictions

```
In [26]: tr_pred4 = np.round(nn4.predict(tr_s[dt_in]).ravel())
         ts_pred4 = np.round(nn4.predict(ts_s[dt_in]).ravel())
```

...And we can evaluate the optimized model:

```
In [27]: tr_c4, tr_f4, tr_s4 = cmodel.cost(tr['machine'].values, tr_pred4, 0.5, return_margin=True)
         ts_c4, ts_f4, ts_s4 = cmodel.cost(ts['machine'].values, ts_pred4, 0.5, return_margin=True)
         print(f'Cost: {tr_c4} (training), {ts_c4} (test)')
         print(f'Avg. fails: {tr_f4/len(tr_mcn)} (training), {ts_f4/len(ts_mcn)} (test)')
         print(f'Avg. slack: {tr_s4/len(tr_mcn):.2f} (training), {ts_s4/len(ts_mcn):.2f} (test)')
```

```
Cost: -18994 (training), -7127 (test)
Avg. fails: 0.0 (training), 0.0 (test)
Avg. slack: 13.95 (training), 10.68 (test)
```

- We are now on-par with the best regression based approaches!
- Actual results may be slightly better or worse due to randomness

Considerations

Why choosing a classification-based approach?

- In some cases, classification can be **easier** than regression
- The threshold affects training, which is computationally heavier
- ...But it means that training can **adapt to the threshold**

We we tackle the prediction and cost optimization stages in isolation:

- The two stages are "glued" by the ML loss (MSE or cross-entropy)
-But the loss does not penalize mistakes based on their (economic) cost!
- Ideally, an end-to-end cost optimization approach should work better

Bayesian/surrogate-based optimization can have many uses

- Hyper-parameter optimization
- Fitting simulation models
- Optimizing equipment setpoints (e.g. Air Conditioning), based on simulation
- ...