

Same Stop Chance

Same Stop Chance

Let's re-examine our RUL-regressor

- We have observed that poor accuracy for high RULs is not much of an issue
- ...But **why is that the case?**

We are going to use the RUL in a condition in the form

$$f(x, \lambda) < \theta$$

- Therefore, what really matters is not the RUL value itself
- ...But the **probability** that the condition is satisfied or violated

This depends entirely on the (lower) quantiles of the distribution of $f(x, \lambda)$

- ...And assumes a probabilistic interpretation for $f(x, \lambda)$
- But is our predictor probabilistic?

Same Stop Chance

Yes, of course! We have an MSE loss, therefore we know that:

$$\operatorname{argmin}_{\lambda} \sum_{i=1}^n (f(\hat{x}_i, \lambda) - \hat{y}_i)^2 = \operatorname{argmax}_{\lambda} \prod_{i=1}^n \phi(f(\hat{x}_i, \lambda) - \hat{y}_i)$$

- I.e. we are implicitly training a regressor with Normally distributed output
- ...Having mean $f(\hat{x}_i, \lambda)$ and **uniform variance** (unary in this case)

As a consequence, we put the same effort in approximating all examples


- But when the RUL is larger, even with a poorer approximation
- ...There will be **still a high chance** that our threshold condition is false

We do **not want to put **the same effort** in all example**

Same Stop Chance

For example, assuming perfect predictions and $\sigma_i = 1 + \hat{y}_i$


```
In [3]: t = np.linspace(0, 100)
y, sigma = t[:-1], 1+t[:-1]
cmapss.plot_rul(pred=y, stddev=sigma, figsize=figsize)
```



Same Stop Chance

For example, assuming perfect predictions and $\sigma_i = 1 + \hat{y}_i$

```
In [3]: t = np.linspace(0, 100)
y, sigma = t[::-1], 1+t[::-1]
cmapss.plot_rul(pred=y, stddev=sigma, figsize=figsize)
```



- The lower quantiles are constant!

Same Stop Chance

So, we can train for this behavior by minimizing:

$$\operatorname{argmax}_{\lambda} \prod_{i=1}^n \frac{1}{\sigma_i} \phi \left(\frac{f(\hat{x}_i, \lambda) - \hat{y}_i}{\sigma_i} \right)$$

And we will choose $\sigma_i = 1 + \hat{y}_i$. By algebraic manipulation we get:

$$\begin{aligned} \operatorname{argmin}_{\lambda} - \sum_{i=1}^n \log \frac{1}{\sigma_i} - \sum_{i=1}^n \log \phi \left(\frac{f(\hat{x}_i, \lambda) - \hat{y}_i}{\sigma_i} \right) &= \\ \operatorname{argmin}_{\lambda} - \sum_{i=1}^n \log \frac{1}{\sigma_i} - \sum_{i=1}^n \log \frac{1}{\sqrt{2\pi}} - \frac{1}{\sigma_i^2} (f(\hat{x}_i, \lambda) - \hat{y}_i)^2 &= \\ \operatorname{argmin}_{\lambda} \sum_{i=1}^n \frac{1}{\sigma_i^2} (f(\hat{x}_i, \lambda) - \hat{y}_i)^2 \end{aligned}$$

Same Stop Chance

We have just established that:

$$\operatorname{argmax}_{\lambda} \prod_{i=1}^n \frac{1}{\sigma_i} \phi \left(\frac{f(\hat{x}_i, \lambda) - \hat{y}_i}{\sigma_i} \right) = \operatorname{argmin}_{\lambda} \sum_{i=1}^n \frac{1}{\sigma_i^2} (f(\hat{x}_i, \lambda) - \hat{y}_i)^2$$

- I.e. training a normally distributed predictor with per-sample variance σ_i^2
- Is equivalent to MSE training with **sample weights** $w_i = 1/\sigma_i^2$

This is suprising simple!

So, let's define the sample weights

```
In [4]: sample_weight = 1/((1+tr_s['rul'].values)**2)
```

- We use the rule $\sigma_i = 1 + \hat{y}_i$ to have identical stop chances

Training

Let's train our MLP architecture with this modification

```
In [5]: nn1 = build_regressor(hidden=[32, 32])
nn1.compile(optimizer='Adam', loss='mse')
history1 = nn1.fit(tr_s[dt_in], tr_s['rul'], validation_split=0.2,
                  sample_weight=sample_weight,
                  callbacks=cb, batch_size=32, epochs=20, verbose=1)
```


```
Epoch 1/20
1135/1135 [=====] - 2s 1ms/step - loss: 0.0090 - val_loss: 0.0069
Epoch 2/20
1135/1135 [=====] - 1s 531us/step - loss: 0.0076 - val_loss: 0.0059
Epoch 3/20
1135/1135 [=====] - 1s 532us/step - loss: 0.0074 - val_loss: 0.0063
Epoch 4/20
1135/1135 [=====] - 1s 536us/step - loss: 0.0072 - val_loss: 0.0065
Epoch 5/20
1135/1135 [=====] - 1s 536us/step - loss: 0.0072 - val_loss: 0.0065
Epoch 6/20
1135/1135 [=====] - 1s 624us/step - loss: 0.0071 - val_loss: 0.0058
Epoch 7/20
1135/1135 [=====] - 1s 627us/step - loss: 0.0071 - val_loss: 0.0056
Epoch 8/20
1135/1135 [=====] - 1s 626us/step - loss: 0.0070 - val_loss: 0.0058
Epoch 9/20
1135/1135 [=====] - 1s 634us/step - loss: 0.0070 - val_loss: 0.0063
Epoch 10/20
1135/1135 [=====] - 1s 628us/step - loss: 0.0070 - val_loss: 0.0058
Epoch 11/20
```


Training

Let's check the loss evolution over time and its final value

```
In [6]: cmapss.plot_training_history(history1, figsize=figsize)
        tr1, vl1 = history1.history["loss"][-1], np.min(history1.history["val_loss"])
        print(f'Loss: {tr1:.4f} (training, final), {vl1:.4f} (validation, best)')
```

Loss: 0.0068 (training, final), 0.0054 (validation, best)




Predictions

The predictions should be (slightly) worse than before in terms of accuracy

```
In [7]: tr_pred1 = nn1.predict(tr_s[dt_in]).ravel() * trmaxrul  
cmapss.plot_pred_scatter(tr_pred1, tr['rul'], figsize=figsize)  
print(f'R2 score: {r2_score(tr["rul"], tr_pred1)}')
```


R2 score: 0.5389392096077179



Predictions

Let's have a look at the predictions over time (on the training set)


```
In [10]: stop = 1095  
cmapss.plot_rul(tr_pred1[:stop], tr['rul'][:stop], figsize=figsize)
```



Predictions

Let's have a look at the predictions over time (on the training set)

```
In [10]: stop = 1095  
cmapss.plot_rul(tr_pred1[:stop], tr['rul'][:stop], figsize=figsize)
```




- Notice how they are noisier in the high-RUL section

Threshold Optimization

Now we need to optimize the threshold as usual

```
In [11]: th_range = np.arange(0, 100)
tr_thr1 = cmapss.opt_threshold_and_plot(tr['machine'].values, tr_pred1, th_range, cmodel, figsize=(10, 5))
print(f'Optimal threshold for the training set: {tr_thr1}')
```

Optimal threshold for the training set: 13



Evaluation

Let's see how it fares in terms of cost

```
In [12]: ts_pred1 = nn1.predict(ts_s[dt_in]).ravel() * trmaxrul
tr_c1, tr_f1, tr_s1 = cmodel.cost(tr['machine'].values, tr_pred1, tr_thr1, return_margin=True)
ts_c1, ts_f1, ts_s1 = cmodel.cost(ts['machine'].values, ts_pred1, tr_thr1, return_margin=True)
print(f'Cost: {tr_c1} (training), {ts_c1} (test)')
print(f'Avg. fails: {tr_f1/len(tr_mcn)} (training), {ts_f1/len(ts_mcn)} (test)')
print(f'Avg. slack: {tr_s1/len(tr_mcn):.2f} (training), {ts_s1/len(ts_mcn):.2f} (test)')
```

```
Cost: -18425 (training), -6819 (test)
Avg. fails: 0.0 (training), 0.0 (test)
Avg. slack: 17.12 (training), 15.57 (test)
```

- The results should be on par than the MLP
 - ...Depending on training randomness
- The number of fails may be slightly larger
 - Forcing equal stop chances tends to lead to slightly more risky models

We did not gain much, but this is a useful trick to know!

Negative-Binomial Model

Negative Binomial Distribution

The negative binomial distribution

- ...Models the probability to have a number "failures"
- ...Before a given number of "successes" is achieved
- ...Assuming a constant success probability p

It's probability mass function is given by:

$$f(k, r, p) = \binom{k + r - 1}{r - 1} (1 - p)^k p^r$$

- k is the number of failures
- r is the number of successes
- The binomial coefficient yields the number of combinations of $r - 1$ successes
- ...Over $k + r - 1$ trials

RUL and NB Distribution

There is a connection between the NB distribution and our process

We can view a "success" as the **end of the run**, a "failure" as an **operating step**

- So, if we assume a constant p for all future steps...
- ...The RUL follows a negative binomial distribution, i.e.:

$$y \sim NB(1, p)$$

- The first distribution parameter is r
- We have $r = 1$ since after a single "success" the run is over

We can use a neural model to estimate p based on the observed data, i.e.:

$$y \sim NB(1, p(\hat{x}_i, \lambda))$$

- By doing so, we effectively obtain **a hybrid neural-probabilistic model**

Training a Neural-Probabilistic Model

We can train our hybrid model for maximum likelihood

Or, better, for minimum negative log likelihood:

$$\operatorname{argmin}_{\lambda} - \sum_{i=1}^n \log f(\hat{y}_i, 1, p(\hat{x}_i, \lambda))$$

- Where f is the probability mass function for the NB distribution
- \hat{y}_i is the RUL value (i.e. the number of "failures"...
- ...1 is the number of successes (end of the run)
- ... $p(\hat{x}_i, \lambda)$ is the "success" probability estimated by the neural model

At inference time:

- The distribution **mean** will provide a RUL estimate
- We will be able to access a **variance**, **quantiles**, and **confidence intervals**!

Building the Probabilistic Model

We will build the probabilistic model using `tensorflow_probability`

In particular, we will build a **custom loss function** for keras:

```
In [13]: import tensorflow_probability as tfp
from tensorflow.keras import backend as k

def negbin_likelihood(y_true, y_pred):
    # y_true = RUL, y_pred = probability of going on
    dist = tfp.distributions.NegativeBinomial(total_count=1, logits=y_pred)
    return -k.sum(dist.log_prob(y_true))
```

- The `NegativeBinomial` class swaps the roles of "success" and "failure"
- In particular, it is designed to work with the "failure" probability (i.e. moving on)
- It supports logit input, in which case a sigmoid is applied to obtain a probability
- ...And it allows the **easy computation of log probabilities**

Training the Hybrid Model

We can use our MLP architecture to estimate the NB logit

```
In [14]: nn2 = build_regressor(hidden=[32, 32])
nn2.compile(optimizer='Adam', loss=negbin_likelihood)
history2 = nn2.fit(tr_s[dt_in], tr['rul'].astype(np.float32), validation_split=0.2,
                  callbacks=cb,
                  batch_size=32, epochs=20, verbose=1)
```

Epoch 1/20

WARNING:tensorflow:@custom_gradient grad_fn has 'variables' in signature, but no ResourceVariables were used on the forward pass.

1135/1135 [=====] - 1s 765us/step - loss: 250.0821 - val_loss: 184.4586

Epoch 2/20

1135/1135 [=====] - 1s 665us/step - loss: 185.5752 - val_loss: 182.1091

Epoch 3/20

1135/1135 [=====] - 1s 662us/step - loss: 183.9809 - val_loss: 181.0170

Epoch 4/20

1135/1135 [=====] - 1s 662us/step - loss: 183.3231 - val_loss: 180.4456

Epoch 5/20

1135/1135 [=====] - 1s 671us/step - loss: 183.0328 - val_loss: 180.0400

Epoch 6/20

1135/1135 [=====] - 1s 667us/step - loss: 182.9550 - val_loss: 179.7785


Epoch 7/20

Training the Hybrid Model

Let's check the loss behavior over time

```
In [15]: cmapss.plot_training_history(history2, figsize=figsize)
tr2, vl2 = history2.history["loss"][-1], np.min(history2.history["val_loss"])
print(f'Loss: {tr2:.4f} (training, final), {vl2:.4f} (validation, best)')
```

Loss: 181.9762 (training, final), 178.6933 (validation, best)



Obtaining the Predictions

The hybrid model prediction corresponds to the **mean** of the NB distribution

First, we need to obtain the estimated probabilities:

```
In [16]: from scipy.special import expit

tr_logits = nn2.predict(tr_s[dt_in]).ravel()
tr_p = expit(tr_logits)

ts_logits = nn2.predict(ts_s[dt_in]).ravel()
ts_p = expit(ts_logits)
```

- `expit` is just the sigmoid function (i.e. reverse of logit)

Obtaining the Predictions

The hybrid model prediction corresponds to the **mean** of the NB distribution

Second, we need to build NB distribution objects

```
In [17]: from scipy.stats import nbinom  
  
tr_dist = nbinom(1, 1-tr_p)  
ts_dist = nbinom(1, 1-ts_p)
```

The scipy NB implementation goes by the textbook definition

- Hence, it is parameterized with the probability of a "success"
- ...i.e. the complement of what tensorflow does

That's why we use $1 - \text{tr_p}$ and $1 - \text{ts_p}$

Obtaining the Predictions

The hybrid model prediction corresponds to the **mean** of the NB distribution

Third, we can obtain the means:

```
In [18]: tr_pred2 = tr_dist.mean()
         ts_pred2 = ts_dist.mean()
```


But we are not limited to that! We can obtain variances, quantiles, etc.

```
In [19]: tr_q1 = tr_dist.ppf(0.25)
         tr_q3 = tr_dist.ppf(0.75)
         ts_q1 = ts_dist.ppf(0.25)
         ts_q3 = ts_dist.ppf(0.75)
```


Predictions

We can plot the predictions **and the 1s/3rd quartiles**

```
In [20]: stop = 1095  
cmapss.plot_rul(tr_pred2[:stop], tr['rul'][:stop], q1_3=(tr_q1[:stop], tr_q3[:stop]), figsize=fi
```




Threshold Optimization

We can perform threshold optimization using **quantiles**

- E.g. using the 1st quartile we have can choose θ so that we stop...
- ...Once the estimated probability of $f(\hat{x}_i, \lambda) \geq \theta$ drops below 25%

```
In [21]: tr_thr2 = cmapss.opt_threshold_and_plot(tr['machine'].values, tr_q1, th_range, cmodel, figsize=f
print(f'Optimal threshold for the training set: {tr_thr2}')
```

Optimal threshold for the training set: 4



Evaluation

Let's see how it fares in terms of cost

```
In [22]: tr_c2, tr_f2, tr_s2 = cmodel.cost(tr['machine'].values, tr_q1, tr_thr2, return_margin=True)
ts_c2, ts_f2, ts_s2 = cmodel.cost(ts['machine'].values, ts_q1, tr_thr2, return_margin=True)
print(f'Cost: {tr_c2} (training), {ts_c2} (test)')
print(f'Avg. fails: {tr_f2/len(tr_mcn)} (training), {ts_f2/len(ts_mcn)} (test)')
print(f'Avg. slack: {tr_s2/len(tr_mcn):.2f} (training), {ts_s2/len(ts_mcn):.2f} (test)')
```

```
Cost: -18372 (training), -6311 (test)
Avg. fails: 0.0 (training), 0.015873015873015872 (test)
Avg. slack: 17.44 (training), 13.13 (test)
```

- The results are again on par with the best approaches
- ...But we have a clearer interpretations and confidence intervals!

Normal Distribution Model

Normal Distribution Model

The NB distribution fits naturally our process, but has a few drawbacks

In particular, the distribution variance is tied to its mean:

$$\text{mean} = \frac{pr}{1-p} \quad \text{variance} = \frac{pr}{(1-p)^2}$$

- We may want to let the model free to adjust its confidence (variance)
- ...Independently on the prediction (mean)

This can be done via a hybrid neural-probabilistic model

- We need a distribution with (at least two) parameters (e.g. the Normal one)
- ...And then we need neural models to estimate both. E.g.:

$$y \sim \mathcal{N}(\mu(\hat{x}, \lambda), \sigma(\hat{x}, \lambda))$$

Building the Architecture

First, we define a function to build the architecture

```
In [24]: def build_probabilistic_regressor(hidden):
    input_shape = (len(dt_in), )
    model_in = keras.Input(shape=input_shape, dtype='float32')
    x = model_in
    for h in hidden:
        x = layers.Dense(h, activation='relu')(x)
    mu_logsig = layers.Dense(2, activation='linear')(x)
    lf = lambda t: tfp.distributions.Normal(loc=t[:, :1], scale=k.exp(t[:, 1:]))
    model_out = tfp.layers.DistributionLambda(lf)(mu_logsig)
    model = keras.Model(model_in, model_out)
    return model
```

- This time we have chosen to rely on the `DistributionLambda`
 - `DistributionLambda` wraps a probability distribution into a layer
- We pass means and (log) standard deviations in a single tensor
 - Therefore the need to unpack them (i.e. `t[:, :1]` and `t[:, 1:]`)

Loss Function

Our loss function will be the negative log likelihood

How can that be implemented?

- At training time, keras repeatedly "calls" the model
- I.e. it literally invokes the `__call__` method
- For `DistributionLambda`, calling returns a distribution object

Therefore we can formulate the log likelihood as follows:

```
In [25]: def dlambda_likelihood(y_true, dist):  
         return -dist.log_prob(y_true)
```

- This is very similar to our previous customized loss
- ...Except that we are returning negative log likelihoods for individual examples
- keras will aggregate by default via a sum, taking into account sample weights
- Our previous formulation couldn't do this

Training

Finally we can train our model

These type of models often needs to be trained for more epochs

```
In [33]: nn3 = build_probabilistic_regressor(hidden=[32, 32])
nn3.compile(optimizer='Adam', loss=dlambda_likelihood)
history3 = nn3.fit(tr_s[dt_in].astype(np.float32), tr_s['rul'].astype(np.float32), validation_s
                  batch_size=32, epochs=30, verbose=1)
```

```
Epoch 1/30
1135/1135 [=====] - 1s 766us/step - loss: -0.5254 - val_loss: -0.7733
Epoch 2/30
1135/1135 [=====] - 1s 683us/step - loss: -0.6871 - val_loss: -0.7841
Epoch 3/30
1135/1135 [=====] - 1s 674us/step - loss: -0.7334 - val_loss: -0.8779
Epoch 4/30
1135/1135 [=====] - 1s 673us/step - loss: -0.7525 - val_loss: -0.8806
Epoch 5/30
1135/1135 [=====] - 1s 673us/step - loss: -0.7785 - val_loss: -0.8649
Epoch 6/30
1135/1135 [=====] - 1s 679us/step - loss: -0.7939 - val_loss: -0.8961
Epoch 7/30
1135/1135 [=====] - 1s 676us/step - loss: -0.8075 - val_loss: -0.7821
Epoch 8/30
1135/1135 [=====] - 1s 679us/step - loss: -0.8132 - val_loss: -0.9175
Epoch 9/30
1135/1135 [=====] - 1s 684us/step - loss: -0.8310 - val_loss: -0.9577
Epoch 10/30
1135/1135 [=====] - 1s 687us/step - loss: -0.8320 - val_loss: -0.9673
```


Obtaining the Predictions

We are interested both in the predicted mean and standard deviation

- Therefore, we cannot simply call `predict`
- Instead, we **call the model** to obtain distribution objects

```
In [34]: tr_prob_pred = nn3(tr_s[dt_in].values)
         ts_prob_pred = nn3(ts_s[dt_in].values)
```


From these, we can obtain means and standard deviations:

```
In [35]: tr_pred3 = tr_prob_pred.mean().numpy().ravel() * trmaxrul
         tr_std3 = tr_prob_pred.stddev().numpy().ravel() * trmaxrul
         ts_pred3 = ts_prob_pred.mean().numpy().ravel() * trmaxrul
         ts_std3 = ts_prob_pred.stddev().numpy().ravel() * trmaxrul
```

Predictions

We can now plot the predictions for the training set


```
In [36]: cmapss.plot_rul(tr_pred3[:stop], tr['rul'][:stop], stddev=tr_std3[:stop], figsize=figsize)
```



Predictions

...And for the test set

```
In [37]: cmapss.plot_rul(ts_pred3[:stop], ts['rul'][:stop], stddev=ts_std3[:stop], figsize=figsize)
```



Evaluation

We can perform threshold optimization and evaluation

We can use either the means or some quantile

```
In [38]: tr_thr3 = cmapss.opt_threshold_and_plot(tr['machine'].values, tr_pred3, th_range, cmodel, plot=True)
print(f'Optimal threshold for the training set: {tr_thr3}')

tr_c3, tr_f3, tr_s3 = cmodel.cost(tr['machine'].values, tr_pred3, tr_thr3, return_margin=True)
ts_c3, ts_f3, ts_s3 = cmodel.cost(ts['machine'].values, ts_pred3, tr_thr3, return_margin=True)
print(f'Cost: {tr_c3} (training), {ts_c3} (test)')
print(f'Avg. fails: {tr_f3/len(tr_mcn)} (training), {ts_f3/len(ts_mcn)} (test)')
print(f'Avg. slack: {tr_s3/len(tr_mcn):.2f} (training), {ts_s3/len(ts_mcn):.2f} (test)')
```

```
Optimal threshold for the training set: 13
Cost: -17915 (training), -6940 (test)
Avg. fails: 0.005376344086021506 (training), 0.0 (test)
Avg. slack: 16.88 (training), 13.65 (test)
```

- The results are once again on par with the best approaches
- As in the previous case, we have a clearer interpretation...
- ...But also more flexible confidence intervals!

Considerations

Why going for hybrid probabilistic models?

- They are typically harder to train than traditional ML models
- ...But they provide confidence **quantiles and standard deviations!**
- **Never** underestimate how useful a confidence interval can be

Probabilistic models enable reasoning

- We can **choose thresholds based on probabilistic considerations**
 - E.g. I want the estimated chance that $RUL \leq 1$ to be lower than 0.10
 - This is **very handy when data is scarce**
 - ...And cost optimization becomes therefore impossible
- It is possible to define probabilistic **cost** models
 - E.g. expected financial cost of maintenance policy
 - They are often the key to build end-to-end cost optimization approaches